

Gaurav Adlakha

Principal / Lead Data Scientist – Generative AI | MLOps | Cloud

Bangalore, India | +91 9717033977 | gaurav.adlakha1509@gmail.com | <https://www.linkedin.com/in/gaurav-adlakha-406b1648/>

PROFESSIONAL SUMMARY

Senior Data Scientist (8 yrs in ML/DL, 16 yrs in software) specializing in large-language-model (LLM) solutions, retrieval-augmented generation (RAG), and end-to-end MLOps on AWS & Kubernetes.

- Cut model deployment time **75 %**, increased information-retrieval relevance **85 %**, and accelerated feature delivery **40 %**.
 - Hands-on with Python, PyTorch, TensorFlow, scikit-learn, MLflow, Kubeflow, Spark, and CI/CD tooling.
-

CORE COMPETENCIES

Machine Learning & AI: Supervised/Unsupervised ML • Ensemble Methods • Recommendation Systems

Deep Learning: CNNs • RNNs • Transformers • LLM Fine-Tuning • GPT-4 • Llama • Hugging Face • LangChain • Prompt Engineering

NLP: RAG • Semantic Search • Embeddings • Text-to-SQL • NER • OCR-enabled NLP

MLOps & DevOps: MLflow • Kubeflow • AWS SageMaker • CI/CD (Git, GitHub Actions, Jenkins) • Docker •

Kubernetes Cloud & Big Data: AWS (Lambda, EC2, EKS, Step Functions, S3) • GCP • Hadoop • Spark • Databricks • Snowflake

Programming: Python • PyTorch • TensorFlow • Keras • scikit-learn • pandas • NumPy • Java • C++

Databases & Storage: PostgreSQL • MongoDB • Vector DBs (FAISS, Pinecone) • Parquet

Visualization & BI: Plotly • Dash • Tableau • Power BI

Leadership: Product Road-mapping • Agile/Scrum • Stakeholder Management • Technical Mentorship

PROFESSIONAL EXPERIENCE

POWERSCHOOL – Principal Data Scientist

Apr 2022 – Present | Bangalore, India

- Architected **scalable RAG pipelines** (FAISS & Pinecone); accuracy ↑ 72 % → 94 %, latency ↓ 300 ms.
- Designed **generative-AI content engine** producing 200+ curriculum items/month, cutting cycle time 60 %.
- Built **teacher-assistant chatbot** (LLM + LangChain) used by 25 K students/day; problem-solving scores ↑ 40 %.
- Implemented **model-safety guardrails** and fine-tuned Llama-2 (Hugging Face); inappropriate outputs ↓ 99.7 %.
- Deployed **MLflow** across 4 teams—deployment window shrank 2 wks → 3 hrs (-85 %).
- Led MLOps roadmap; CI/CD with **GitHub Actions** boosted release velocity 40 %.

PwC Labs – Senior Data Scientist (Advisory)

Sep 2020 – Apr 2022 | Bangalore, India

- Spearheaded NLP/OCR platform processing **3 M+** docs/yr at 85 + % accuracy, cutting manual effort 70 %.
- Authored **config-driven Python framework** (pandas, scikit-learn, FastAPI); dev cycle time -40 %.
- Orchestrated distributed training on **Kubernetes (15 nodes)**; 2 TB data, training time -65 %.
- Delivered **skill-recommendation engine** saving \$1.5 M/yr and improving mobility 45 %.

Mphasis – Data Scientist

Mar 2019 – Sep 2020 | Bangalore, India

- Built FedEx classifier (**92 %** accuracy, 200 K shipments/day, 1 K labels).
- Deployed DL invoice extractor; accuracy ↑ 75 % → 93 %, processing time ↓ 20 %.
- Rolled out **serverless AWS Lambda** stack, cutting infra cost 89 % (\$12 K → \$1.3 K/mo).

Accenture – Data Scientist

Oct 2016 – Jan 2019

- Automated seismic-data pipelines for 15 TB datasets; analysis time ↓ 70 %.
- Developed ML algorithms for geological detection (**93 %** accuracy) enabling 3 new drill sites.

Earlier technical roles (IBM, CSG, Nucleus, etc.) available on request.

EDUCATION

- **B.E. Computer Science**, M.D. University, 2008
- **Deep Learning Specialization**, Deeplearning.ai, 2019